CS365 Operating System Midterm Exam

Name: _____

Question Number	Score
1	
2	
3	
4	
5	
6	
7	
Total:	

Question1:
Explain the following terminology: (1pt each)

1. **System Call**
2. **Process**
3. **I/O-bound Process**
4. **Context Switch**
5. **LWP**
6. **Symmetric Multiprocessing**

1. **System Call:** is how a program requests service from an operating system's kernel
2. **Process:** a program in execution.
3. **I/O-bound process:** spends more time doing I/O than computations, many short CPU bursts.
4. **Context Switch:** is the process of storing and restoring the state (the context) of a process or thread so that execution can be resumed from the same point at a later time.
5. **LWP:** is an intermediate data structure between user and kernel threads, which appears to be a virtual processor on which process can schedule user thread to run on a kernel thread.
6. **Symmetric Multiprocessing:** In the symmetric multiprocessor system with two or more identical processors, each processor is self-scheduling and each processor may have its own ready queue.

Question2:

2.1 Explain and describe 3 different types of process schedulers. (3 pts)

- **Long-term scheduler**
(or job scheduler) –selects which processes should be brought into the ready queue

- **Short-term scheduler**
(or CPU scheduler) –selects which process should be executed next and allocates CPU

- **Midterm Scheduler**

Sometimes it can be advantage to remove process from memory and thus decrease the degree of multiprogrammimg. This scheme is called s**wapping**

2.2 Describe 3 different multithreading models for mapping user threads to kernel threads. Also, describe at least **one** advantage or disadvantage for each model. (5 pts)

Many to one: It maps many user-level threads to one kernel thread. The entire system may block makes a block system call.

One to one: Each user-level thread has one corresponding kernel thread. The only drawback is that creating a user thread requires creating the corresponding kernel thread

Many to many: Multiplexes many user-level threads to a smaller or equal number of kernel threads. User can create as many threads as they want

2.3 Draw the state diagram of a process from its creation to termination, including all transitions, and briefly elaborate every state and every transition . (5 pts)

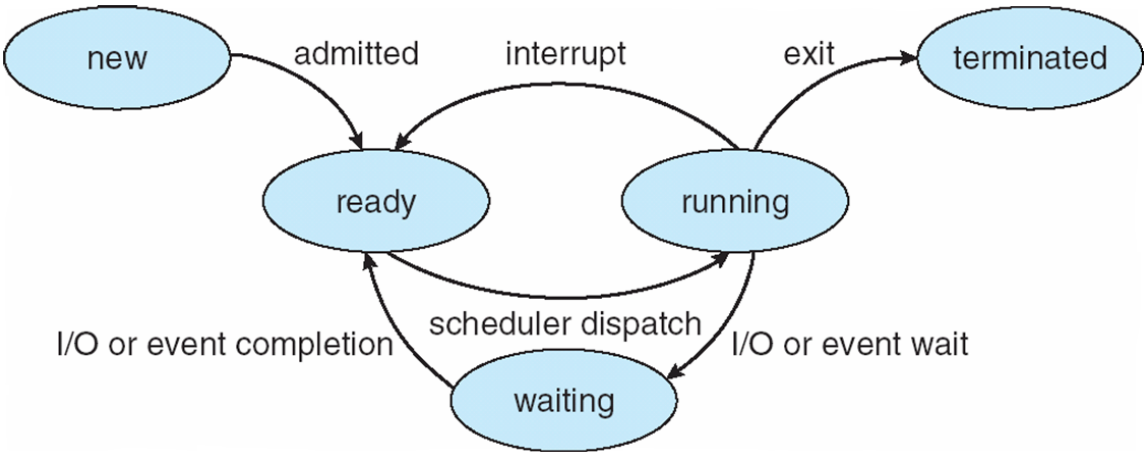

Question3:
According to the bounded buffer code in shared memory system we discussed in the class, draw the ring structure with size 4 and indicate the location of "in" and "out" for each time space. (If the buffer is full or empty, indicates how many jobs are waiting) (10pts, 2 pts for each)

Producer View

```
while (true) {
   /* produce an item and put in next Produced*/
   while (count == BUFFER_SIZE)
        ; // do nothing

   buffer [in] = nextProduced;
   in = (in + 1) % BUFFER_SIZE;
   count++;
}
```

Consumer View

```
while (true) {
   while (count == 0)
        ; // do nothing

   nextConsumed= buffer[out];
   out = (out + 1) % BUFFER_SIZE;
   count--;
   /* consume the item in next Consumed
}
```

Time	Action	waiting:0	count:0	in:1	out:1
1	Producer produces 2 jobs	0	2	3	1
2	Consumer consumes 1 job	0	1	3	2
3	Producer produces 3 jobs	0	4	2	2
4	Consumer consumes 3 jobs	0	1	2	1
5	Consumer consumes 3 jobs	1	0	2	2

Question4:

By "**Multilevel Queue Fixed priority**" scheduling algorithm, draw the CPU scheduling Gantt chart (5 pts) and complete the table for the given processes information. (10 pts)

	Process	Burst time	Arriving time	Algorithm
■ 1st Foreground	P1	50	0.0	RR interval:20
	P2	15	30.0	(RR is a non-preemptive
	P3	45	30.0	algorithm)
■ 2nd Foreground	P4	40	0.0	SJF Preemptive
	P5	10	120.0	
■ Background	P6	30	60.0	FCFS
	P7	20	130.0	

Gantt Chart:

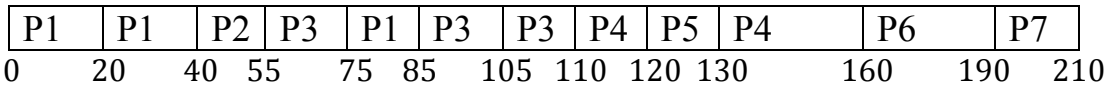

Scheduling criteria:

	P1	P2	P3	P4	P5	P6	P7
Waiting time	35	10	35	120	0	100	60
Turnaround time	85	25	80	160	10	130	80
Response time	0	10	25	110	0	100	60

Question5:

Circle the best answer. (14 pts, 1 pt for each)
AAACBDBBBDDCDB
1. The major difficulty in designing a layered operating system approach is ___.

A) appropriately defining the various layers

B) making sure that each layer hides certain data structures, hardware, and operations from higher-level layers
C) debugging a particular layer
D) making sure each layer is easily converted to modules

2. A boot block ____.
A) typically only knows the location and length of the rest of the bootstrap program
B) typically is sophisticated enough to load the operating system and begin its execution
C) is composed of multiple disk blocks
D) is composed of multiple disk cylinders

3. Microkernels use _____ for communication.
A) message passing
B) shared memory
C) system calls
D) virtualization

4. The UNIX fork() system call creates a new process. What is the equivalent system call in WINDOWS:
A) NTCreateProcess()
B) process()
C) CreateProcess ()
D) getpid()

5. All access to POSIX shared memory requires a system call.
A) True
B) False

6. When a child process is created, which of the following is a possibility in terms of the execution or address space of the child process?
A) The child process runs concurrently with the parent.
B) The child process has a new program loaded into it.
C) The child is a duplicate of the parent.
D) All of the above

7. Cancellation points are associated with ____ cancellation.
A) asynchronous
B) deferred
C) synchronous
D) immediate

8. The _____ model allows a user-level thread to be bound to one kernel thread.
A) many-to-many
B) two-level
C) one-to-one
D) many-to-one

9. Which of the following is true of cooperative scheduling?
A) It requires a timer.
B) A process keeps the CPU until it releases the CPU either by terminating or by switching to the waiting state.
C) It incurs a cost associated with access to shared data.
D) A process switches from the running state to the ready state when an interrupt occurs.

10. ____ is the number of processes that are completed per time unit.
A) CPU utilization
B) Response time
C) Turnaround time
D) Throughput

11. ____ scheduling is approximated by predicting the next CPU burst with an exponential average of the measured lengths of previous CPU bursts.
A) Multilevel queue
B) RR
C) FCFS
D) SJF

12. The ____ scheduling algorithm is designed especially for time-sharing systems.
A) SJF
B) FCFS
C) RR

13. Which of the following scheduling algorithms must be nonpreemptive?
A) SJF
B) RR
C) FCFS
D) priority algorithms

14. Which of the following is true of multilevel queue scheduling?
A) Processes can move between queues.
B) Each queue has its own scheduling algorithm.
C) A queue cannot have absolute priority over lower-priority queues.
D) It is the most general CPU-scheduling algorithm.